

Problem Set 1

AUTHOR

Taeyoung Heo(UMID # 2914 5801)

GitHub repository link : https://github.com/tedheo10/homework1_STATS506

Problem 1 - Wine data

P1-a. import "wine.data" into a data.frame in R and give column names

```
winedata <- read.table("wine.data", sep = ",")
colnames(winedata) <- c("class", "alcohol", "acid", "ash", "alcalinity",
"magnesium", "phenols", "flavanoids", "nonflavanoid_phenols",
"proanthocyanins", "color_intensity", "hue", "protein", "proline")
```

P1-b. check the number of wines within each class

[class1 = 59, class2 = 71, class3 = 48]

```
wineclass <- winedata$class
count <- c(length(which(wineclass==1)), length(which(wineclass == 2)),
           length(which(wineclass == 3)))
report <- c(59, 71, 48)

if(count[1] == report[1] & count[2] == report[2] & count[3] == report[3])
{
  print("correct as reported in wine.names")
} else {
  print("wrong as reported in wine.names")
}
```

[1] "correct as reported in wine.names"

P1-c-1. the correlation between alcohol content and color intensity

the correlation : **0.5463642**

```
alcohol <- winedata$alcohol
color <- winedata$color_intensity
cor(alcohol, color)
```

[1] 0.5463642

P1-c-2. the highest and the lowest correlation

the highest : **class1(0.4082913)**

the lowest : **class2(0.2697891)**

```

class <- winedata$class

alcohol1 <- winedata$alcohol[winedata$class == 1]
alcohol2 <- winedata$alcohol[winedata$class == 2]
alcohol3 <- winedata$alcohol[winedata$class == 3]

color1 <- winedata$color[winedata$class == 1]
color2 <- winedata$color[winedata$class == 2]
color3 <- winedata$color[winedata$class == 3]

gcor <- c(cor(alcohol1, color1), cor(alcohol2, color2),
          cor(alcohol3, color3))
gcor

```

```
[1] 0.4082913 0.2697891 0.3503777
```

P1-c-3. the alcohol content of the wine with the highest color intensity
the alcohol content with the highest color intensity : **14.34**

```
length(which(color == max(color)))
```

```
[1] 1
```

```
alcohol_hcolor <- winedata$alcohol[winedata$color_intensity == max(color)]
alcohol_hcolor
```

```
[1] 14.34
```

P1-c-4. the percentage of the wines with a higher content of proanthocyanins compare to ash
the percentage : **8.426966%**

```

proanthocyanins <- winedata$proanthocyanins
proan_higher <- winedata$proanthocyanins[winedata$proanthocyanins > winedata$ash]
length(proan_higher) / length(proanthocyanins) * 100

```

```
[1] 8.426966
```

P1-d. table for the average value of each variable
the table : **average_table**

```

average_table <- matrix(0, nrow = 4, ncol = length(colnames(winedata))-1)
colnames(average_table) <- colnames(winedata[2:14])
rownames(average_table) <- c("overall", "class1", "class2", "class3")

cal_mat <- as.matrix(winedata)
n <- ncol(cal_mat)-1

ave_overall <- numeric(n)

```

```

ave_class1 <- numeric(n)
ave_class2 <- numeric(n)
ave_class3 <- numeric(n)

for ( i in 1:n) {
  ave_overall[i] <- mean(cal_mat[,i+1])
  ave_class1[i] <- mean(cal_mat[cal_mat[,1] == 1, i+1])
  ave_class2[i] <- mean(cal_mat[cal_mat[,1] == 2, i+1])
  ave_class3[i] <- mean(cal_mat[cal_mat[,1] == 3, i+1])
}

average_table["overall",] <- ave_overall
average_table["class1",] <- ave_class1
average_table["class2",] <- ave_class2
average_table["class3",] <- ave_class3

average_table

```

	alcohol	acid	ash	alkalinity	magnesium	phenols	flavanoids
overall	13.00062	2.336348	2.366517	19.49494	99.74157	2.295112	2.0292697
class1	13.74475	2.010678	2.455593	17.03729	106.33898	2.840169	2.9823729
class2	12.27873	1.932676	2.244789	20.23803	94.54930	2.258873	2.0808451
class3	13.15375	3.333750	2.437083	21.41667	99.31250	1.678750	0.7814583

	nonflavanoid_phenols	proanthocyanins	color_intensity	hue	protein
overall	0.3618539	1.590899	5.058090	0.9574494	2.611685
class1	0.2900000	1.899322	5.528305	1.0620339	3.157797
class2	0.3636620	1.630282	3.086620	1.0562817	2.785352
class3	0.4475000	1.153542	7.396250	0.6827083	1.683542

	proline
overall	746.8933
class1	1115.7119
class2	519.5070
class3	629.8958

P1-e. t-test to examine whether the level of phenols differs across the three classes

We can reject the hypothesis that the level of phenols is the same across the three classes because all three p-values are less than 0.01. So we can conclude that the level of phenols differs across the three classes.

t-test p-value between class1 and class2 : **1.889e⁻¹¹**

t-test p-value between class1 and class3 : **< 2.2e⁻¹⁶**

t-test p-value between class2 and class3 : **1.622e⁻¹⁰**

```

phenols1 <- winedata$phenols[winedata$class == 1]
phenols2 <- winedata$phenols[winedata$class == 2]
phenols3 <- winedata$phenols[winedata$class == 3]

t.test(phenols1, phenols2)

```

```
data: phenols1 and phenols2
t = 7.4206, df = 119.14, p-value = 1.889e-11
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 0.4261870 0.7364055
sample estimates:
mean of x mean of y
 2.840169  2.258873
```

```
t.test(phenols1, phenols3)
```

Welch Two Sample t-test

```
data: phenols1 and phenols3
t = 17.12, df = 98.356, p-value < 2.2e-16
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 1.026801 1.296038
sample estimates:
mean of x mean of y
 2.840169  1.678750
```

```
t.test(phenols2, phenols3)
```

Welch Two Sample t-test

```
data: phenols2 and phenols3
t = 7.0125, df = 116.91, p-value = 1.622e-10
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 0.4162855 0.7439610
sample estimates:
mean of x mean of y
 2.258873  1.678750
```

Problem 2 - AskAManager.org Data

P2-a. import AskAManager.csv into a data.frame

```
salary <- read.csv("AskAManager.csv")
```

P2-b. simplify the variable names

```
colnames(salary) <- c("number", "time", "age", "industry", "title",
                      "title_context", "salary", "bonus", "currency",
                      "currency_other", "income_context", "country",
```

```
"state", "city", "experience", "experience_field",  
"education", "gender", "race")
```

P2-c. restrict the data to those paid in USD

the number of observations before : **28,062**

the number of observations after the restriction : **23,374**

```
length(salary$currency)
```

```
[1] 28062
```

```
length(salary$currency[salary$currency == "USD"])
```

```
[1] 23374
```

```
salary_res <- salary[salary$currency == "USD", ]  
  
nrow(salary_res) == length(salary$currency[salary$currency == "USD"])
```

```
[1] TRUE
```

P2-d. Eliminate any rows for which their age, years of experience in their field, and years of experience total are impossible

the number of observation after the elimination : **23,116**

```
age <- salary_res$age  
field <- salary_res$experience_field  
total <- salary_res$experience  
  
age[!duplicated(age)]
```

```
[1] "25-34"      "45-54"      "35-44"      "18-24"      "65 or over"  
[6] "55-64"      "under 18"
```

```
field[!duplicated(field)]
```

```
[1] "5-7 years"      "2 - 4 years"      "21 - 30 years"      "11 - 20 years"  
[5] "8 - 10 years"    "1 year or less"    "31 - 40 years"      "41 years or more"
```

```
total[!duplicated(total)]
```

```
[1] "5-7 years"      "2 - 4 years"      "8 - 10 years"      "21 - 30 years"  
[5] "11 - 20 years"    "41 years or more" "31 - 40 years"      "1 year or less"
```

```
age_kind <- c("25-34" = 16, "45-54" = 36, "35-44" = 26, "18-24" = 6,  
             "65 or over" = 56, "55-64" = 46, "under 18" = 0)  
  
experience_kind <- c("5-7 years" = 5, "2 - 4 years" = 2,
```

```

      "8 - 10 years" = 8, "21 - 30 years" = 21,
      "11 - 20 years" = 11, "41 years or more" = 41,
      "31 - 40 years" = 31, "1 year or less" = 1 )

```

```

age_trans <- age_kind[age]
field_trans <- experience_kind[field]
total_trans <- experience_kind[total]
n_usd <- nrow(salary_res)
memo <- logical(n_usd)

for(i in 1:n_usd) {
  # if age < 18, eliminate the data
  # if age - 18 < field experience, eliminate the data
  # if age - 18 < total experience, eliminate the data
  # if total experience < field experience, eliminate the data

  if(age_trans[i] == 0) {
    memo[i] = FALSE
  }
  else if(field_trans[i]>age_trans[i]) {
    memo[i] = FALSE
  }
  else if(total_trans[i]>age_trans[i]) {
    memo[i] = FALSE
  }
  else if(field_trans[i]>total_trans[i]) {
    memo[i] = FALSE
  }
  else {
    memo[i] = TRUE
  }
}

#test whether age is under 18
salary_year <- salary_res[memo, ]
salary_year$age[!duplicated(salary_year$age)]

```

```

[1] "25-34"      "45-54"      "35-44"      "18-24"      "65 or over"
[6] "55-64"

```

```

#test example for whether their is reasonable with total experience
experience_test <- salary_year$age[salary_year$experience
                                == "41 years or more"]
experience_test[!duplicated(experience_test)]

```

```

[1] "65 or over" "55-64"

```

```

#test example for whether their is reasonable with field experience
fieldexp_test <- salary_year$age[salary_year$experience_field
                                == "41 years or more"]
fieldexp_test[!duplicated(fieldexp_test)]

```

```

[1] "65 or over" "55-64"

```

```
#test example for whether filed experience is less than total experience
exp_fieldepx_test <- salary_year$experience_field[salary_year$experience
                                                    == "8 - 10 years"]

exp_fieldepx_test[!duplicated(exp_fieldepx_test)]
```

```
[1] "5-7 years"      "2 - 4 years"    "8 - 10 years"   "1 year or less"
```

```
nrow(salary_year)
```

```
[1] 23116
```

P2-e. eliminate any rows with extremely low or high salaries.

The other monetary compensation is added to salary in order to eliminate false salaries.

The extreme high salary : **1,459,608 USD** (mean + 2*standard deviation)

-> There are values on the right side of the distribution that differ from the mean by more than twice the standard deviation.

The extreme low salary : **1,000 USD**

-> There's a very thick tail on the left side of the distribution.

The sample size : **23,044**

```
income <- salary_year$salary
bonus <- salary_year$bonus
bonus[is.na(bonus)] <- 0
monetary_reward <- income + bonus
summary(monetary_reward)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0	58000	81000	106240	120000	102000000

```
n_year <- length(monetary_reward)

mean_monetary <- mean(monetary_reward)
sd_monetary <- sd(monetary_reward)
mean_monetary + 2*sd_monetary
```

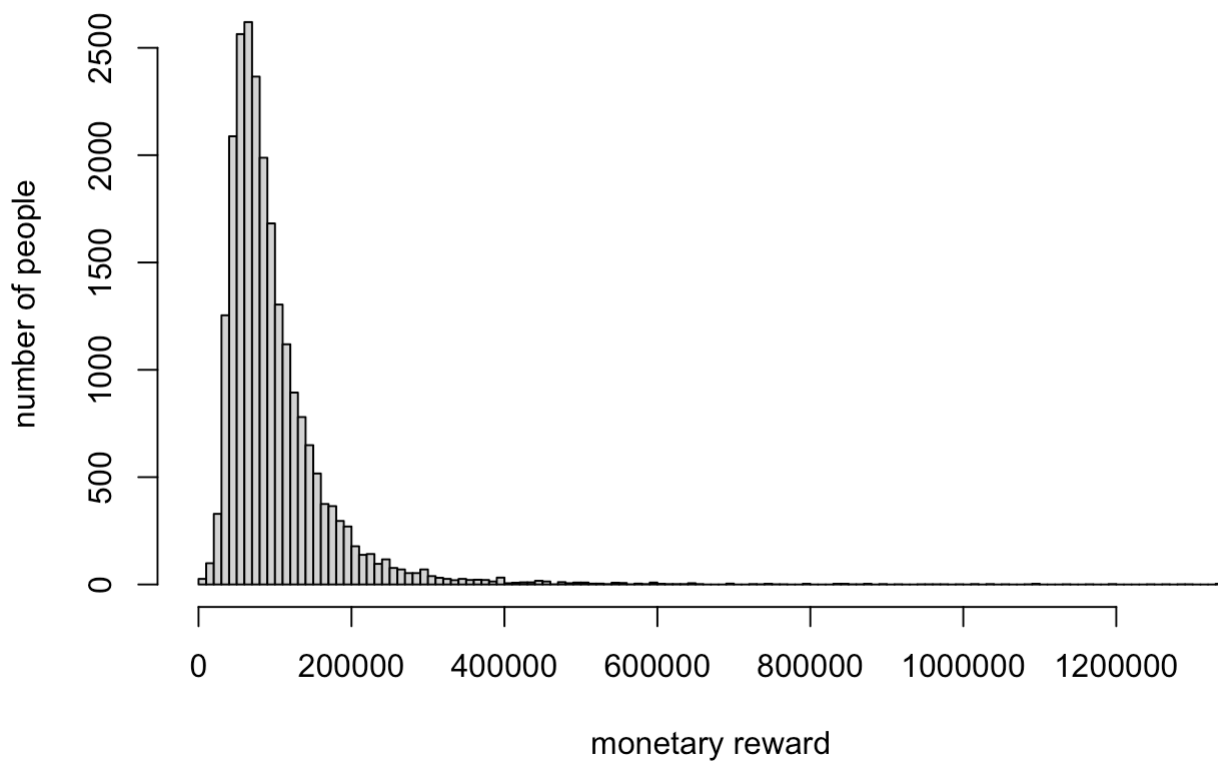
```
[1] 1459608
```

```
mean_monetary - sd_monetary
```

```
[1] -570443.5
```


```
hist(monetary_reward[monetary_reward <= 1459608 & monetary_reward
                    >= 1000], breaks = 100, main =
     "Monetary reward between 1,000 and 1,459,608 USD",
     xlab = "monetary reward", ylab = "number of people")
```

Monetary reward between 1,000 and 1,459,608 USD



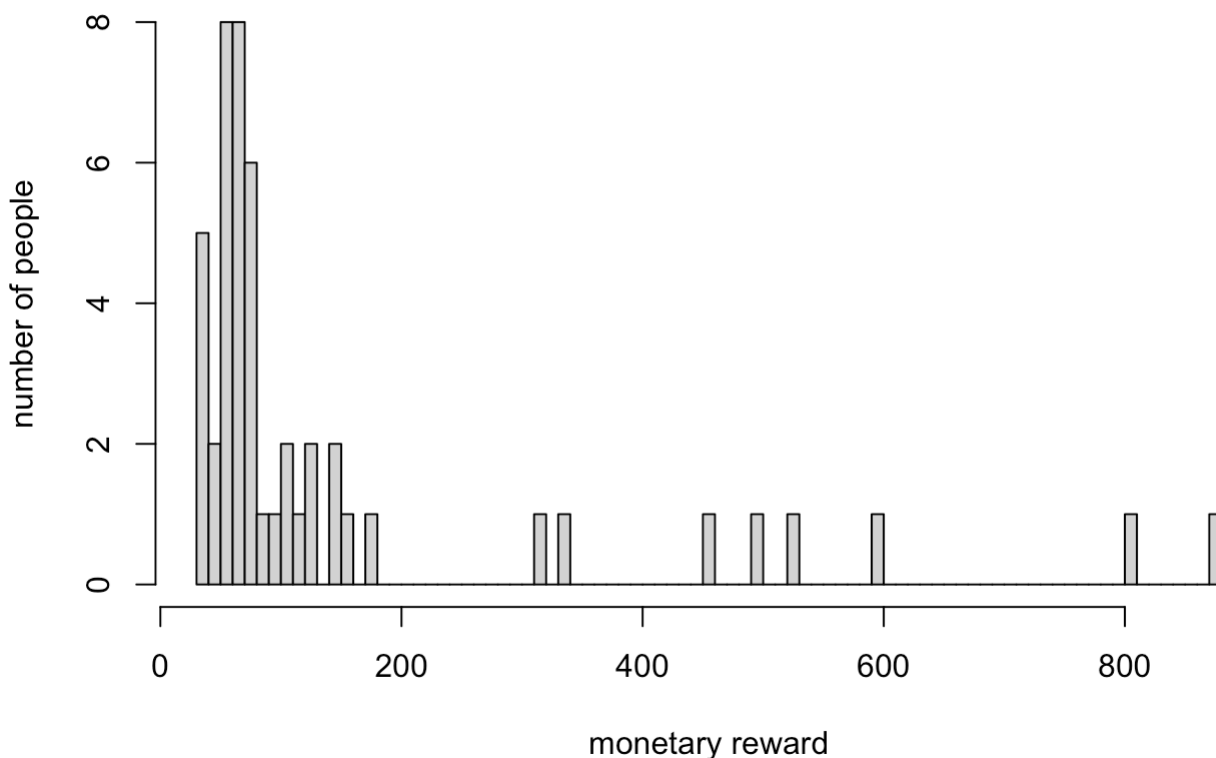
```
hist(monetary_reward[monetary_reward > 1459608], breaks = 100, main =  
      "Monetary reward over 1,459,608 USD",  
      xlab = "monetary reward", ylab = "number of people")
```


Monetary reward over 1,459,608 USD



```
hist(monetary_reward[monetary_reward < 1000 & monetary_reward > 0],  
     breaks = 100, main = "Monetary reward less than 1,000 USD",  
     xlab = "monetary reward", ylab = "number of people")
```

Monetary reward less than 1,000 USD



```
salary_year$monetary_reward <- monetary_reward
salary_monetary <- salary_year[monetary_reward <= 1459608 & monetary_reward >=
nrow(salary_monetary)
```

[1] 23044

Problem 3 - Palindromic Numbers

P3-a. function "isPalindromic" that check if a given positive integer is a palindromic

```
#' Palindromic of Integers
#'
#' Determine whether a positive intger is palindromic or not
#'
#' This function takes a positive integer x and returns a logical value
#' whether the integer is a palindromic number or not and return reversed #'
#' number.
#'
#' @param x a positive integer
#' @param na.rm logical, should NA's be removed? Default is FALSE
#' @return list of a logical value and a numeric value
isPalindromic <- function(x, na.rm = FALSE) {
  if(!is.numeric(x)) {
    # this code is from STATS 506 class note
    warning("x must be a positive integer, attempting to convert")
```

```

suppressWarnings(x <- as.numeric(x))
if(all(is.na(x))) {
  stop(
    "x must be a positive integer or convertible to a positive integer")
}
}
if(length(x) == 0) {
  # this code is from STATS 506 class note
  stop("x must have strictly positive length")
}
if(!is.logical(na.rm)) {
  # this code is from STATS 506 class note
  warning("na.rm must be logical")
}
if(x <= 0 ) {
  stop(
    "x must be a positive integer or convertible to a positive integer")
}
if(x != round(x)) {
  stop(
    "x must be a positive integer or convertible to a positive integer")
}

y <- as.character(x)
z <- as.numeric(unlist(strsplit(y, "")))
k <- length(z)
check_palindromic <- logical(k)
reverse <- numeric(k)
for(i in 1:k) {
  if(z[i] == z[k+1-i]) {
    check_palindromic[i] <- TRUE
  } else {
    check_palindromic[i] <- FALSE
  }
  reverse[i] <- z[k+1-i]
}
reverse_return <- as.numeric(paste(reverse, collapse = ""))
if(all(check_palindromic == TRUE)) {
  return(list(isPalindromic = TRUE, reversed = reverse_return))
}
else {
  return(list(isPalindromic = FALSE, reversed = reverse_return))
}
}

isPalindromic(1100) # check the "isPalindromic()" function

```

```

$isPalindromic
[1] FALSE

```

```

$reversed
[1] 11

```

P3-b. create a nextPalindrome function that finds the next palindromic number strictly greater than the input

```
#' Next Palindrome
#'
#' Find the next palindromic number strictly greater than the input
#'
#' This function takes a positive integer x and check from x+1 to the next #'
#' and return the next palindromic number.
#'
#' @param x a positive integer
#' @param na.rm logical, should NA's be removed? Default is FALSE
#' @return a numeric value
nextPalindrome <- function(x, na.rm = FALSE) {
  if(!is.numeric(x)) {
    # this code is from STATS 506 class note
    warning("x must be a positive integer, attempting to convert")
    suppressWarnings(x <- as.numeric(x))
    if(all(is.na(x))) {
      stop(
        "x must be a positive integer or convertible to a positive integer")
    }
  }
  if(length(x) == 0) {
    # this code is from STATS 506 class note
    stop("x must have strictly positive length")
  }
  if(!is.logical(na.rm)) {
    # this code is from STATS 506 class note
    warning("na.rm must be logical")
  }
  if(x <= 0) {
    stop(
      "x must be a positive integer or convertible to a positive integer")
  }
  if(x != round(x)) {
    stop(
      "x must be a positive integer or convertible to a positive integer")
  }

  y <- as.character(x+1)
  z <- as.numeric(unlist(strsplit(y, "")))
  k <- length(z)
  check <- logical(k)
  next_palindrome <- numeric(k)
  while(!all(check == TRUE)) {
    for(i in 1:k) {
      if(z[i] == z[k+1-i]) {
        check[i] <- TRUE
      } else {
        check[i] <- FALSE
      }
    }
  }
```

```

        next_palindrome[i] <- z[i]
      }
      y <- as.character(as.numeric(y)+1)
      z <- as.numeric(unlist(strsplit(y, "")))
      k <- length(z)
    }
    next_palindrome <- as.numeric(paste(next_palindrome, collapse = ""))
    return(nextPalindrome = next_palindrome)
  }

  nextPalindrome(1523) # check the "nextPalindrome()" function

```

[1] 1551

P3-c. Use above functions to find the next palindrome for each of the following: (391, 9928, 19272719, 109, 2)

i. 391 : **393**

ii. 9928 : **9999**

iii. 19272719 : **19277291**

iv. 109 : **111**

v. 2 : **3**

```
nextPalindrome(391)
```

[1] 393

```
nextPalindrome(9928)
```

[1] 9999

```
nextPalindrome(19272719)
```

[1] 19277291

```
nextPalindrome(109)
```

[1] 111

```
nextPalindrome(2)
```

[1] 3